

Report: Object and Sub-Object Detection System

Github Repository: <https://github.com/jayxdev/sub-obj-detector>

1. Introduction

The objective of this project was to develop a robust computer vision system capable of detecting objects and their associated sub-objects in a hierarchical structure. The system is designed to detect various objects (e.g., “person”, “car”) and their corresponding sub-objects (e.g., “helmet”, “tire”) in video streams. The solution must provide accurate and adaptable detection while maintaining real-time performance on a CPU. It also includes a functionality for retrieving cropped images of specific sub-objects for further analysis.

2. System Overview

The computer vision system leverages a fine-tuned YOLOv8n model to perform real-time object and sub-object detection. The YOLO model is trained to recognize different objects, and boundaries are used to identify sub-objects within those main objects. The system processes video frames in real-time and outputs detection results in a specified hierarchical JSON format. Furthermore, it supports the retrieval of sub-object images by cropping the identified regions from the video frames.

3. Object and Sub-Object Detection

- **Main Object Detection:** The system detects various objects such as “person”, “car”, “dog”, etc., in each frame of the video stream.
- **Sub-Object Detection:** Once the main objects are detected, sub-objects (e.g., head for persons, tires for cars, persons inside bus) are identified within the bounding boxes of the parent objects.
- **Hierarchical Structure:** The sub-objects are linked to their parent objects based on their relative positions. The system uses bounding box relationships to associate sub-objects with parent objects, creating a hierarchy where each object is uniquely indexed.

Example:

- For a detected person, the system will check for sub-objects such as “helmet” within the person's bounding box. The sub-object's bounding box will be captured, and the hierarchical relationship will be saved.

4. JSON Output Format

The detection results are structured in the following hierarchical JSON format:

```
```json
{
 "object": "Person",
 "id": 1,
 "bbox": [x1, y1, x2, y2],
 "subobject": {
 "object": "Helmet",
 "id": 1,
 "bbox": [x1, y1, x2, y2]
 }
}
```

- **object:** Name of the main object.
- **id:** Unique identifier for the detected object.
- **bbox:** Bounding box coordinates for the object in the format [x1, y1, x2, y2].
- **subobject:** A nested dictionary for each sub-object associated with the parent object, including the sub-object's name, ID, and bounding box.

## 5. Sub-Object Image Retrieval

- The system allows users to retrieve cropped images of specific sub-objects. For example, after detecting a person, users can extract an image of the helmet (sub-object) associated with that person.
- The system ensures accurate image retrieval by referencing the hierarchical relationships between objects and sub-objects. The bounding box coordinates are used to crop the frame and save sub-object images.

**Example:** For the "Person" object, the system can crop and save the image of the "Helmet" sub-object if it is detected.

## 6. Inference Speed Optimization

- **Optimization:** The system is optimized to process video inputs at 10–30 frames per second (FPS) on a CPU. This ensures that the system meets real-time performance requirements while balancing detection accuracy.
- **Benchmarking:** The system was tested on a sample video, and the inference speed was measured. The average FPS achieved was 20 FPS, which satisfies the performance requirements for real-time video processing.

### Optimization Strategies:

- Model size and inference speed were improved by fine-tuning the YOLOv8n model, optimizing the bounding box determination for sub-object detection, and using batch processing for video frames.
- The system utilizes efficient frame resizing techniques to reduce the computational load while maintaining detection accuracy.

## 7. Modularity and Extensibility

The design of the system is modular, allowing for the easy addition of new object-sub-object pairs (e.g., adding new classes such as "Bike" and "Handlebars"). This modular approach ensures that the system can be extended to handle different object detection scenarios.

- **Instructions for Adaptation:** To adapt the system to other object detection scenarios, users can update the object detection classes and the corresponding sub-objects in the configuration files. The detection model can be retrained with new classes if necessary.

## 8. Challenges and Limitations

- **Occlusion Handling:** The system handles occlusion scenarios by ensuring the object and sub-object detections are refined with non-maximal suppression to filter overlapping bounding boxes. However, in cases of severe occlusion, detection accuracy may decrease.

- **Object Sub-Object Relationships:** The accuracy of sub-object detection largely depends on the robustness of the YOLO model and the clarity of the bounding boxes. Any discrepancies in bounding box determination can affect the sub-object extraction process.

## 9. Benchmarking Results

- **Test Video:** A test video of traffic scenes was used to benchmark the performance of the system.

- **Inference Speed:** On a CPU, the system achieved an average of 20 FPS for the test video, meeting the performance requirement of 10–30 FPS.

- **Detection Accuracy:** The system demonstrated high detection accuracy with an average confidence score of 0.5 for main objects and sub-objects.

## 10. Conclusion

The system successfully meets the objectives of object and sub-object detection with hierarchical relationships, real-time video processing, and sub-object image retrieval. The modular design ensures that the system can be easily extended to handle additional object-sub-object pairs. With performance optimizations in place, the system operates efficiently on a CPU, ensuring scalability for edge devices.

---

## Deliverables

**1. Codebase:** The full source code for the object detection system is available in the provided GitHub repository.

**2. JSON Outputs:** Sample JSON outputs have been provided for the test video, demonstrating the hierarchical structure of detected objects and sub-objects.

**3. Benchmarking Report:** The benchmarking results, including FPS and detection accuracy, are documented above.

**4. Demo:** An optional demo video showcasing real-time object and sub-object detection is included. Detected Objects are marked green and sub-object in blue.